

Ansh Singhal

anshsinghal3107@gmail.com
anshsinghal.dev
github.com/AnshSinghal

Greater Noida, India
+91-8929554991
linkedin.com/in/ansh-singhal

SUMMARY

AI/ML Engineer building high-performance, end-to-end machine learning systems through robust systems design, efficient architecture, scalable deployment pipelines, and backend integration.

SKILLS

- **Programming Languages:** Python, Java, C++, SQL
- **Machine Learning & AI:** PyTorch, TensorFlow, Hugging Face Transformers, LangChain, LlamaIndex
- **Backend & Web Frameworks:** Spring Boot, Flask, FastAPI, Django, REST APIs, gRPC, Bootstrap
- **DevOps & MLOps Tools:** Docker, Kubernetes, Redis, Kafka, Celery, Git, Prometheus, Grafana, CI/CD
- **Cloud Platforms:** Google Cloud Platform (GCP), Microsoft Azure, Amazon Web Services (AWS)
- **Systems Design:** API Design, Observability, Scalable Microservices, Distributed Systems

EXPERIENCE

- **Open Source Contributor – IntelOwl** Remote
Threat Intelligence Platform Jan 2025 – May 2025
 - **Threat Analyzer Development:** Built modular file-based and observable-event analyzers to extract actionable threat intelligence from files, URLs, and system behaviors.
 - **Containerization & APIs:** Containerized analyzers using **Docker** and exposed them via **FastAPI** for reproducible, scalable, and isolated runtime execution.
 - **Django Backend:** Developed and maintained the **Django** backend to expose RESTful endpoints for OSINT enrichment, enabling seamless integration.
 - **Documentation & Dev Onboarding:** Authored detailed documentation and test suites to streamline contributor onboarding and increase adoption in the open-source community.

PROJECTS

- **Agentic RAG – Legal Research Engine:** Built an **Agentic RAG pipeline** for domain-specific legal retrieval with **97% accuracy** and **P95 latency < 500ms**. Designed multi-agent routing via LangGraph and hybrid retrieval using Pinecone (dense + sparse).
Highlights: 2-stage reranking (BGE + MiniLM), LegalSeg-based NLP, and containerized FastAPI microservices deployed on AWS ECS.
Tech: Python, FastAPI, LangGraph, Pinecone, Docker, AWS ECS
- **Flood-GAN – Physics-Aware Flood Mapping using Multi-Channel Conditional GANs:** Designed a **physics-informed cGAN** for **SAR→Optical translation** with flood detection. Achieved **PSNR 31.25**, **SSIM 0.94**, and **L1 0.05**. Model outputs RGB + Water maps with hydrology-consistent supervision.
Highlights: NDWI-consistent RGB generation, cloud-masked losses, dual-head U-Net, PatchGAN discriminator, and speckle preservation loss.
Tech: PyTorch, Lightning, Sentinel-1/2, VGG19, EMA, TTUR
- **Talent Insider – Job Platform:** Developed a scalable recruitment system supporting **1K+ users** with secure bcrypt-based authentication and live CRUD job management. Engineered RESTful APIs and optimized DB operations for sub-100ms queries.
Highlights: Modular MVC structure, asynchronous API calls, and relational data consistency using MySQL transactions.
Tech: Java, Spring Boot, JDBC, MySQL, Maven, Java Swing

EDUCATION

- **Bennett University** Uttar Pradesh, India
B.Tech in Computer Science (AI Specialization); CGPA: 9.65 2023 – 2027

HONORS AND AWARDS

- 6th Rank, Smart India Hackathon – National Inter-College Level
- 1st Rank, International English Olympiad & International Science Olympiad
- Winner, State-Level Debate Championship
- Winner, National Science Exhibition

VOLUNTEER EXPERIENCE

- **Youth Advisory Council - Rotary International (USA)** Remote
Board Member (Only member from Asia selected) Jul 2023 – Jun 2024
 - Collaborated with a team of leaders to shape the future of Rotary’s youth exchange programs while maintaining strong communication and collaboration with other members and staff.
 - Selected as the only member from Asia based on proven leadership in Rotary and its programs, contributing valuable experience and knowledge to guide future youth exchange initiatives.